

Automated Machine Learning

Frank Zijun Zhang, PhD

Division of AI in medicine, Cedars-Sinai Medical Center

UCLA QCBio Collaboratory

AutoML workshop, Winter 2023

Workshop Overview

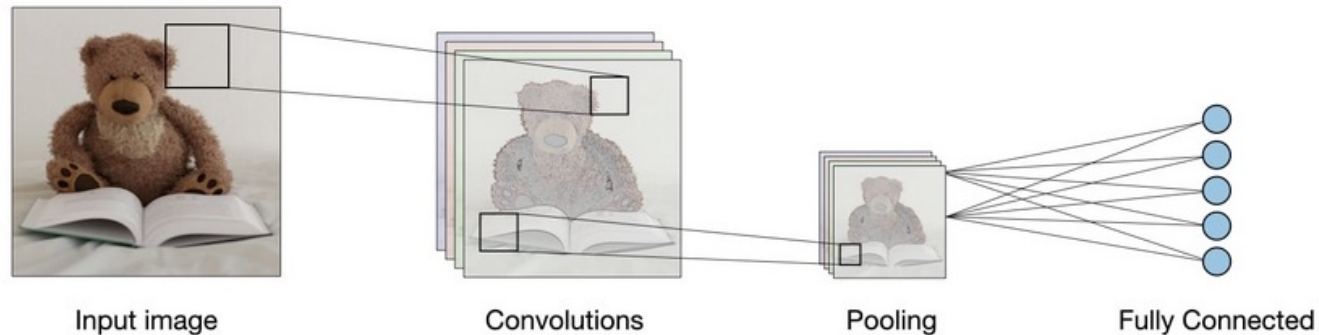
- This is a two-day applied course for building convolutional neural networks and AutoML in genomics.
 - applied: focus on implementations and coding, with a high-level theoretical understanding
 - CNNs: deep learning models that are good at learning spatial patterns (spatial inductive bias)
 - AutoML: automate the process of deep learning design
 - genomics: input DNA sequences, output biological insights

Workshop Goals

- Understand how CNNs model genomic sequences
- Understand the basic concepts in AutoML
- Implement CNNs in PyTorch and Tensorflow
- Implement AutoML prototype in AMBER
- Implement basic model explanation techniques

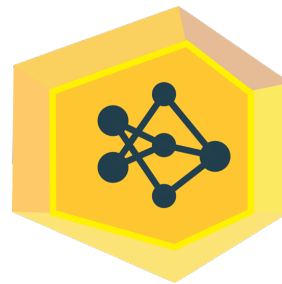

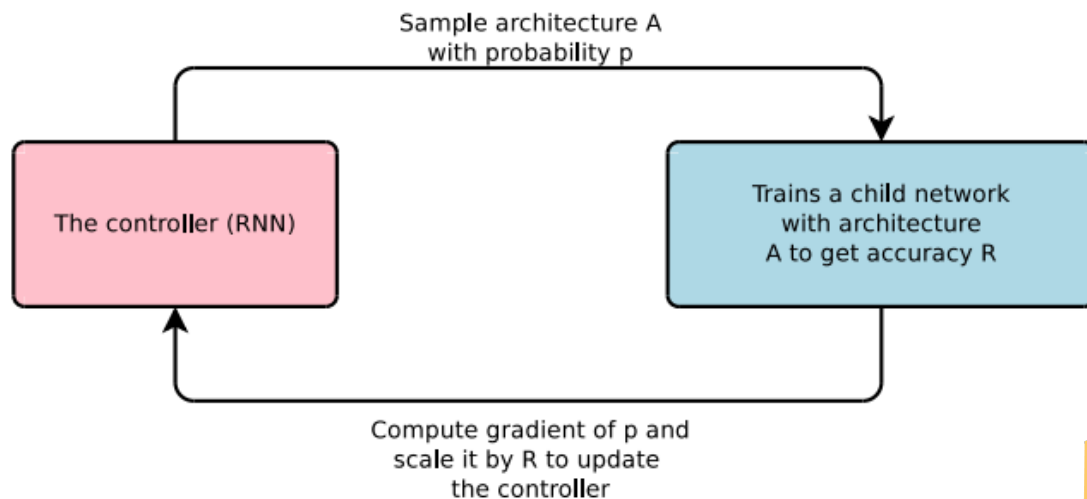
Agenda

- Day1: Build CNNs and Model Evaluation
 - Understand how CNNs model genomic sequences
 - Build CNNs with Tensorflow and PyTorch
 - Model evaluations and tuning




Agenda

- Day2: AutoML and Model Explanation
 - Introduction of reinforcement learning
 - Apply NAS to automate CNN tuning
 - Interpret model and sequence motifs



AMBER

A note on Google Colab and Github

- Go to Github Repo:
- <https://github.com/zhanglab-aim/ucla-automl-workshop>
- Open “UCLA_AutoML_workshop_Day1.ipynb”.
- Click on the Icon  Open in Colab
- Each practice is ~10min; basic code is already implemented in the Notebook
- Explore the Practice questions shown in the slides.

Day 1: Build CNNs and Model Evaluation

Frank Zijun Zhang, PhD


Division of AI in medicine, Cedars-Sinai Medical Center

UCLA QCBio Collaboratory

AutoML workshop, Winter 2023


Recap: machine learning

- ML gives computers the ability to learn without being explicitly programmed.. *but how?*
- Machine learning originated from statistical learning.



$$P(Y = 1) = \text{sigmoid}(\beta_0 + \beta_1 \times \text{petal} + \beta_2 \times \text{sepal})$$

betas are optimized by maximum likelihood.




Recap: machine learning

- Learning Betas with gradient descent
- Optimization: maximize the likelihood on the observed feature-label pairs
- Iteratively update beta to increase the model's likelihood, for each feature-label pair (x_i, y_i) :
 - $\hat{\beta}_t = \hat{\beta}_{t-1} + \eta \frac{\partial}{\partial \beta} \log p(y_i | x_i; \beta)$
 - η : learning rate, controls how big each update is

Recap: machine learning

- Statistics vs machine learning vs deep learning (aka artificial intelligence)
- They are closely connected, but have developed important, specialized subfields

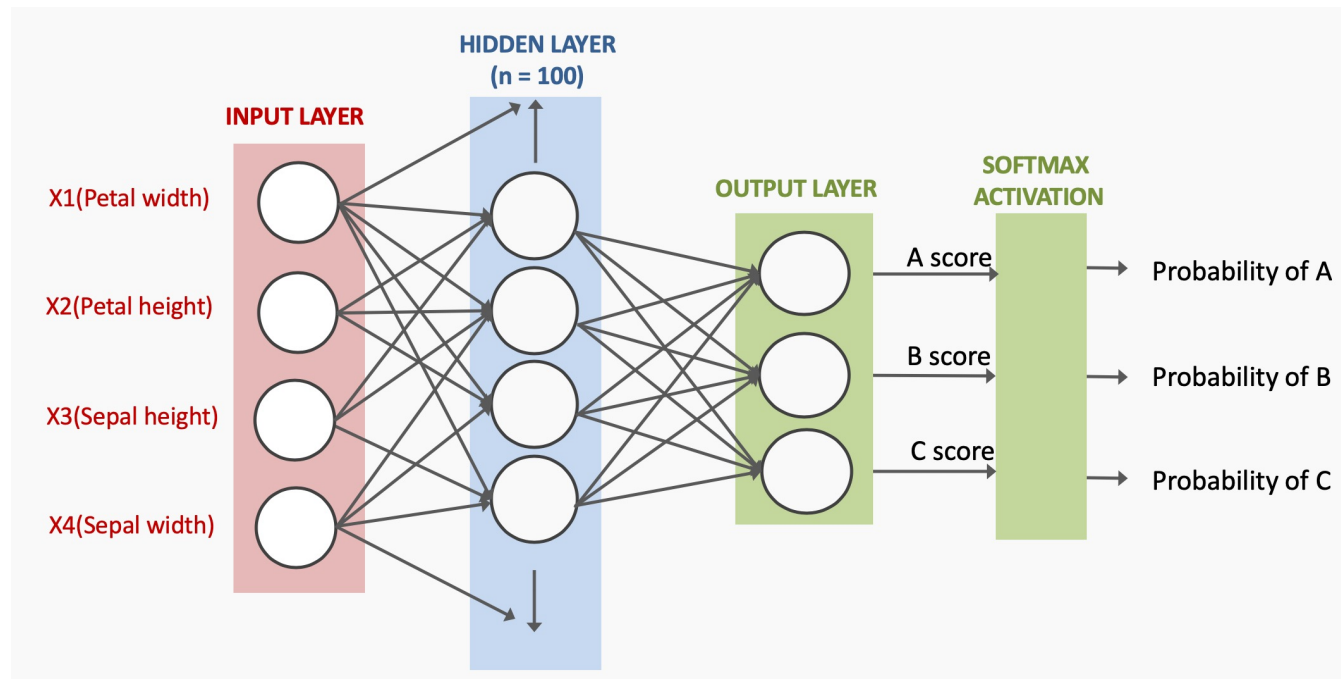


original comic by [sandserif](#)



Deep Neural Networks: Multi-layer perceptron (MLP)

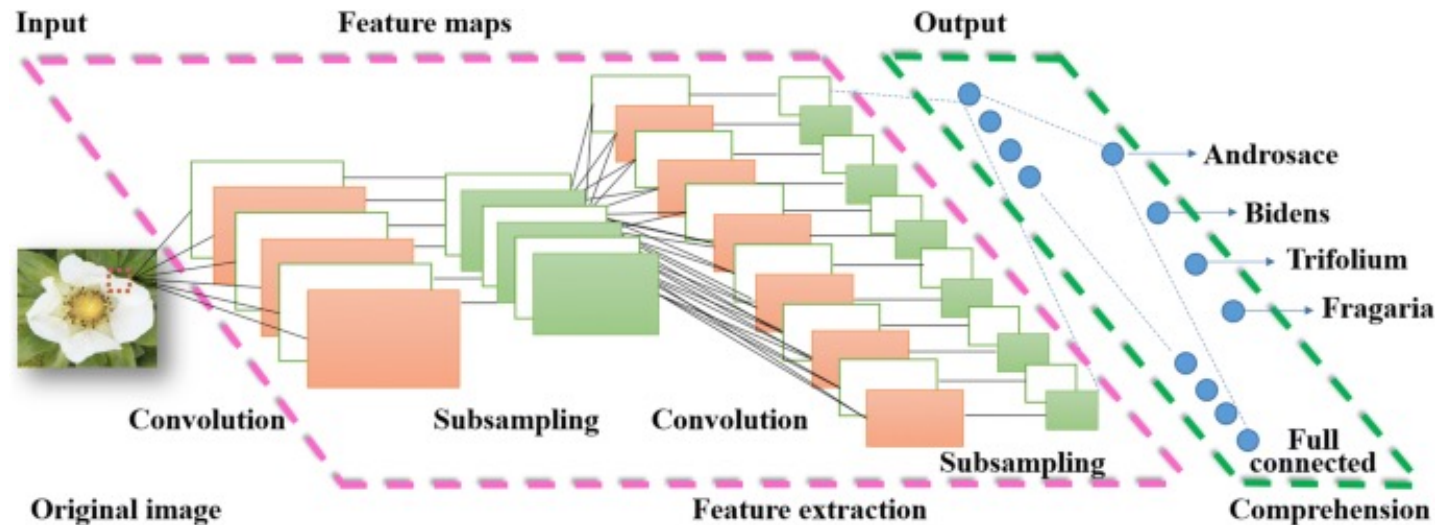
- In its most basic form, DNNs are stacks of logistic regression.



Each one neuron of $n=100$ in the hidden layer is a logistic regression!

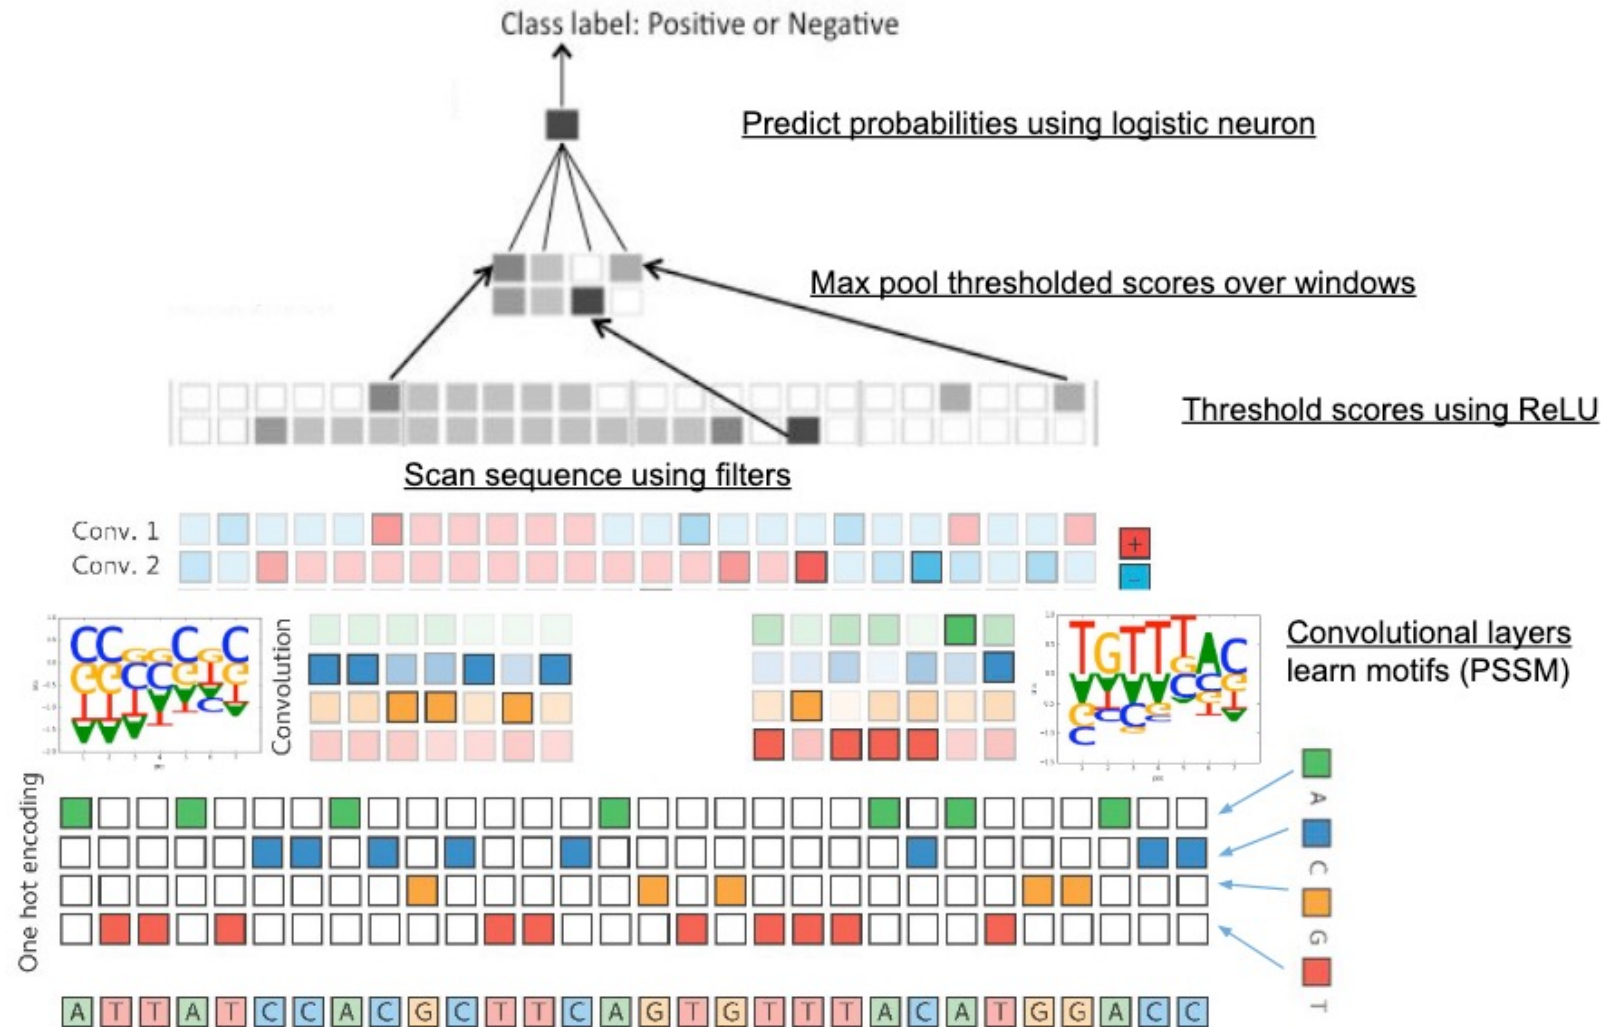
Convolutional Neural Networks on Images

- CNNs are special DNNs that are good at learning local spatial patterns.
- These patterns eliminate the need for manual features (e.g. petal length).
- However, Exact interpretation of image patterns is an active research direction.
 - In contrast, CNN interpretations on genomics are more straightforward.



Deep convolutional neural network for automatic discrimination between *Fragaria* × *Ananassa* flowers and other similar white wild flowers in fields. Plant Methods, 2018.

Convolutional Neural Networks on Genomics




How to train your DragonNN.

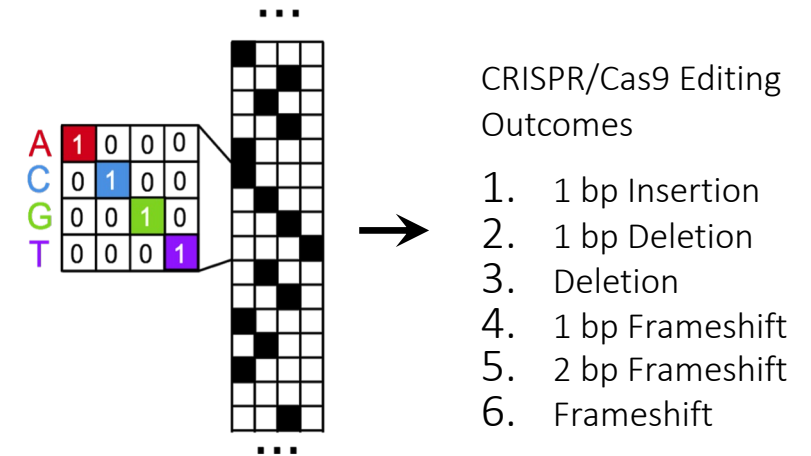

https://www.genome.gov/sites/default/files/Multimedia/Slides/ENCODE2016-ResearchAppsUsers/Anshul_slides.pdf

Effective models of various molecular variations

- CNNs were first developed to predict epigenetic markers from DNA sequences
- Later introduced to vast different molecular variations, such as enhancers, gene expression, RNA splicing/polyadenylation.



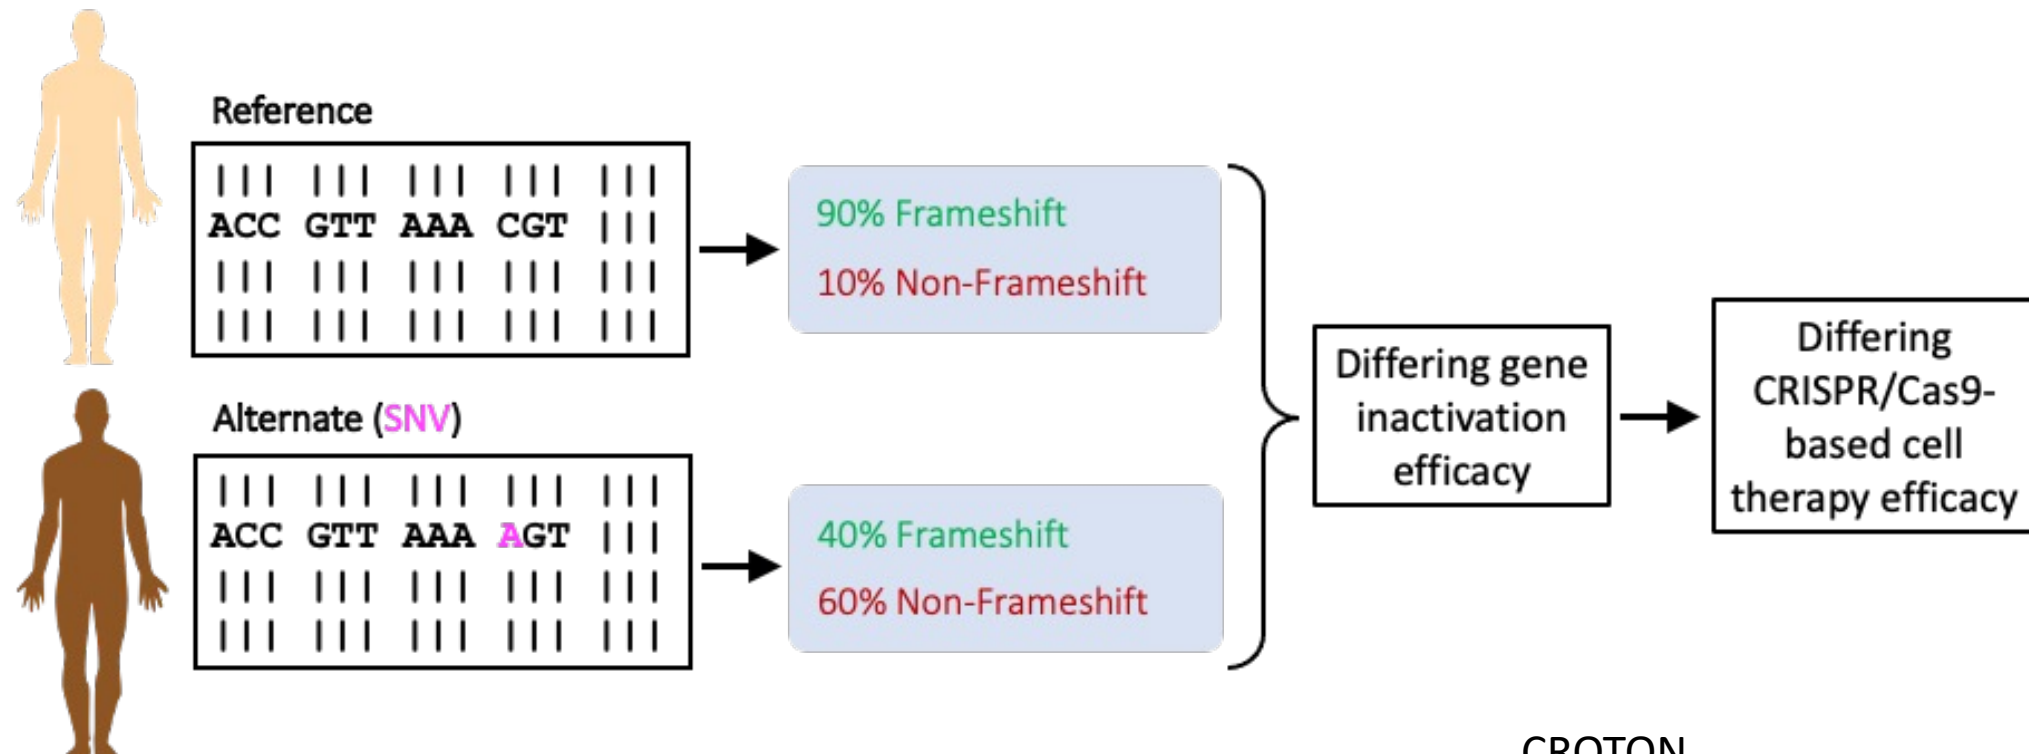
Sei
Chen et al., Nature Genetics, 2022



CROTON
Li et al., Bioinformatics, 2022

Interpretations of functional genetic variations

- Central dogma established the causality direction.
- With a CNN model, we can introduce mutations *in silico* and predict its molecular effect.



CROTON

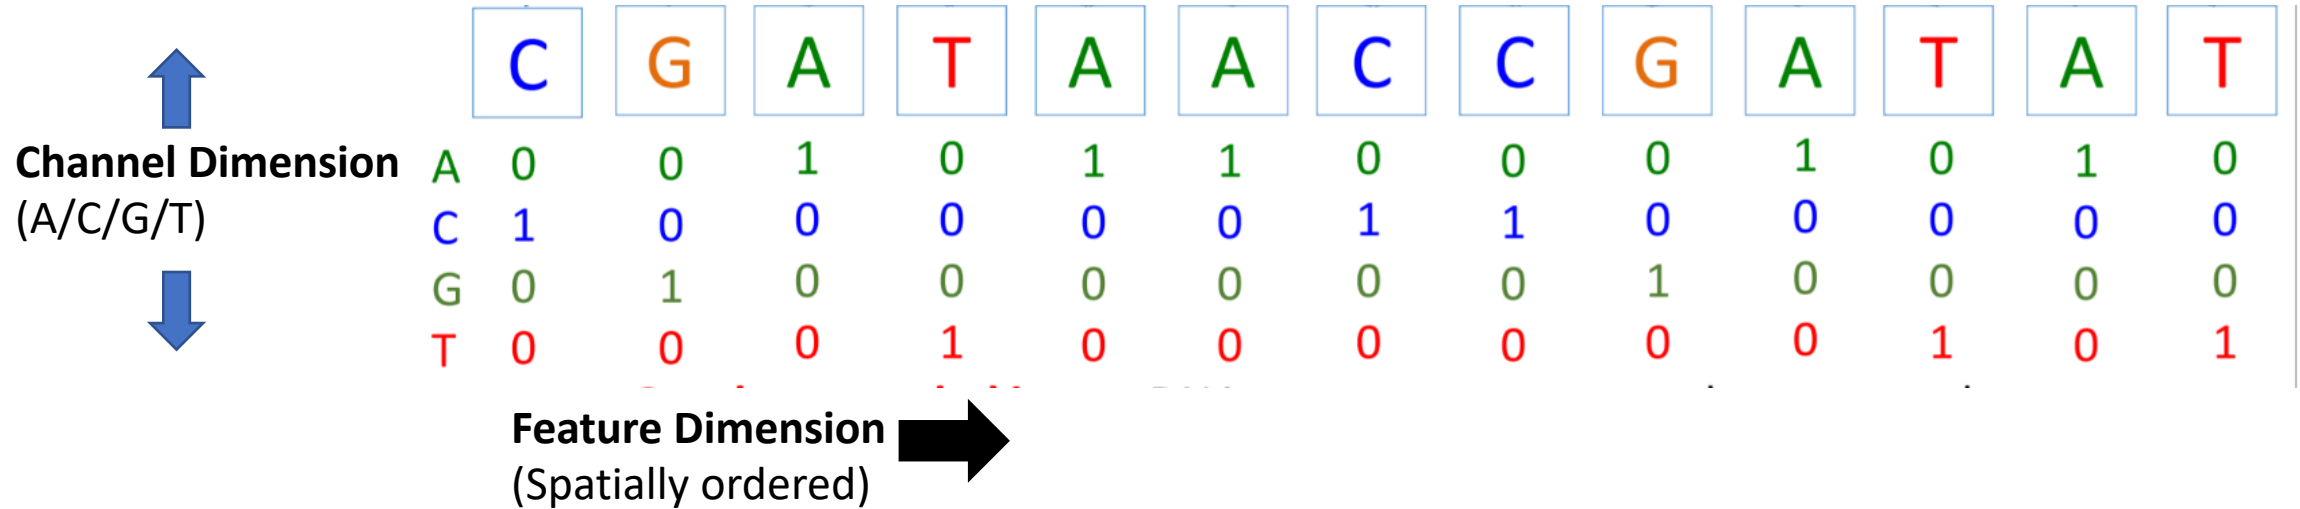
Li et al., Bioinformatics, 2022

CNNs provide useful tools in genomics

- In summary, CNNs are powerful tools in genomics for these reasons:
 - end-to-end modeling without manual feature engineering
 - flexibility to model multiple relevant tasks
 - ability to introduce *in silico* mutations
- Next, we will study and implement a CNN in a step-by-step guide.

Genomic Sequence Input

- Convert DNA sequences to one-hot encoded matrix



Train-Validation-Test Split

- We held-out a random portion of our labelled data to evaluate how well the model will generalize to new cases.
- For many ML methods and DL in particular, an additional validation dataset is used to assess overfitting during training and for tuning hyperparameters.
- Be cautious about potential data leakage-
 - In real human data, it is best practice to hold-out some chromosomes.




Illustration of dataset split.

[Source](#).

Practice 1: Prepare Datasets


- Go to Github Repo and find Day 1 Jupyter notebook
- Click on “Open in Colab”

 Open in Colab

Day 1. Building CNNs to Model Genomic Sequences

Summary:

- In Google Colab, connect to a new Runtime
- Run from the beginning till the end of Section 1. Load Data and Train-Test Split




Practice 1: Prepare Datasets

- 1. Exploratory data analysis on the positive and negative label proportions in training/validation/test sets.
- 2. Exploratory data analysis on the one-hot encoded feature matrices – proportions of A/C/G/T.
- 3. Write your own Python function to convert one-hot matrix back to DNA sequence, and the inverse function that convert DNA sequence to one-hot matrix.

Convolution

- A convolutional kernel is a matrix that can be translated to a motif positional weight matrix.
- The weights are learned through gradient descent (similar to logistic regression's betas).




kernel_size=3

Convolution

- Convolution operation applies a kernel matrix to scan the input sequence.
- Output dim = input feature dim – kernel size + 1 = 5 – 3 + 1 = 3


2.5



Convolution

- Convolution operation applies a kernel matrix to scan the input sequence.
- Output dim = input feature dim – kernel size + 1 = 5 – 3 + 1 = 3

2.5 0.0




Convolution

- Convolution operation applies a kernel matrix to scan the input sequence.
- Output dim = input feature dim – kernel size + 1 = 5 – 3 + 1 = 3

2.5 0.0 -1.5

Convolution kernel

Input




Convolution

- Convolution operation applies a kernel matrix to scan the input sequence.
- Output dim = input feature dim – kernel size + 1 = 5 – 3 + 1 = 3


2.5 0.0 -1.5

Convolution kernel




Convolution

- A convolution layer is a collection of convolution kernels.
 - Input: feature dim=5 (5bp), channel dim=4 (A/C/G/T)
 - Conv: filters=2, kernel size=3
 - Output: feature dim=3, channel dim=2




Pooling

- Scan the input feature map with a sliding window on the feature dimension, and keep the max or average values within each window.
- Pooling changes the feature dim, but not channel dim.




Pooling

- Scan the input feature map with a sliding window on the feature dimension, and keep the max or average values within each window.
- Pooling changes the feature dim, but not channel dim.



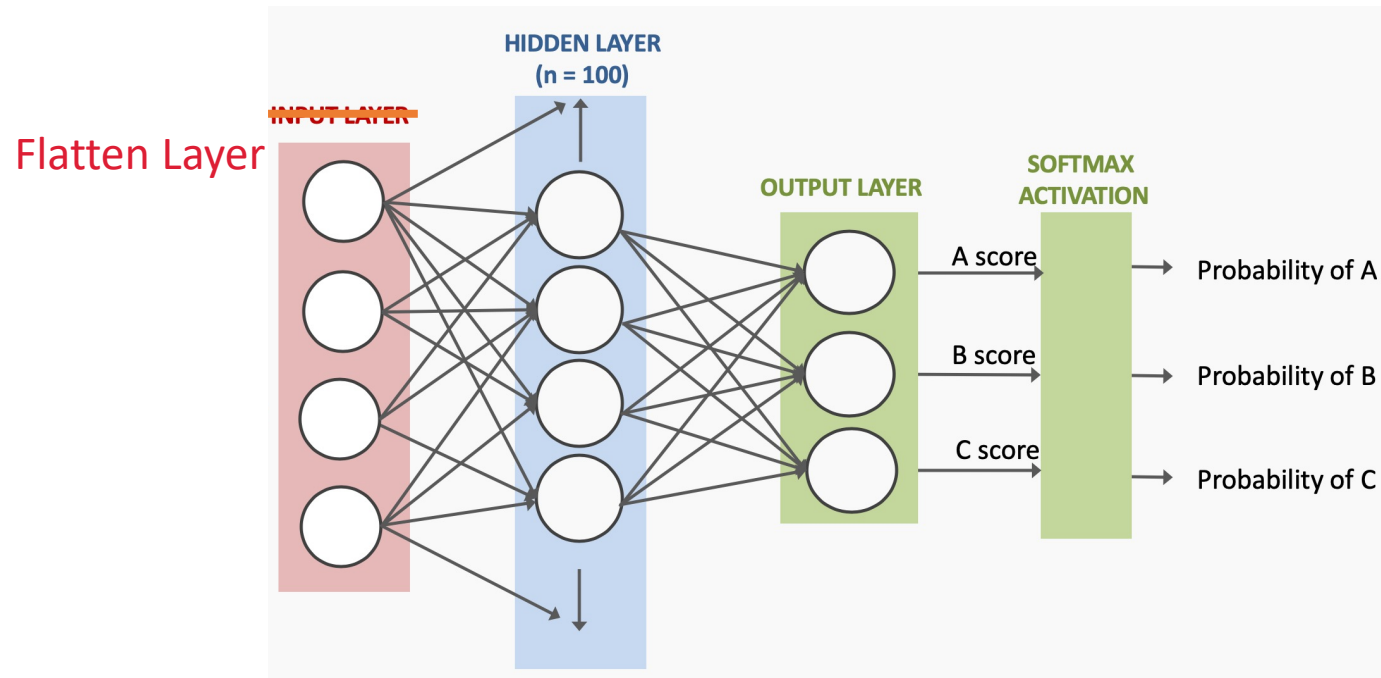
Flatten

- Flatten operations will reshape the feature map from a matrix to a vector.



Fully Connected


- The previous flatten layer brings us back to the multi-layer perceptron.



Why Keras?

- Keras is excellent for fast prototyping and research development.
- The Model API implemented many useful callbacks and metrics with a self-contained design.
- Keras is flexible for advanced models through custom layers.

Model groups by framework



 **Kipoi:** Model zoo for genomics

<https://kipoi.org/>

Practice 2: Build CNN in Keras

Run Section 2.Build a CNN Model using Tensorflow/Keras



- Change the parameter values of `filters`, `kernel_size`, `pool_size`, and observe the shape changes in `model.summary()`.
- Change the flatten layer from GAP to Flatten.
- Uncomment Tensorboard to visualize training results.


Understanding Keras Model API

- `model.compile()`: configure optimizations
- *optimizer*: gradient descent and its variants; Adam is usually a good default choice.
- *loss*: loss function; must be differentiable to parameters.
 - classification: `binary_crossentropy`, `categorical_crossentropy`
 - regression: `mse`, `mae`
- *metrics*: additional monitors for model performance, such as AUROC (will cover next).

Understanding Keras Model API

- `model.fit()`: train model parameters
- *batch_size*: compute gradients on this number of samples
- *epochs*: how many epochs to train; a complete iteration through all train data is one epoch.
- *callbacks*: a list of callback functions that will be triggered upon epoch end.
 - `TensorBoard`: log model performance to tensorboard
 - `ModelCheckpoint`: save model parameters
 - `EarlyStopping`: stop training earlier than specified epochs, if model starts overfitting

Comparing PyTorch vs Keras

- PyTorch is another popular deep learning framework.
- Excellent for research and development of novel models, with more exposure and control over the model's behavior.
- Run Section 3.Build a CNN using Pytorch/Lightning.  [Open in Colab](#)

Building CNNs in PyTorch

- In general, PyTorch model has a slightly higher requirement Python Programming:
 - Models are subclassing `torch.nn.Module`
 - Operators and Tensors are built during `__init__()`
 - Output channel dims are determined by users
 - `model.forward()` sets the forward pass
 - train loops, metrics are custom-built, or relies on other packages (e.g. lightning, torchmetrics)

Can you match the PyTorch code implementations to Keras syntax?


Model Evaluations

- Evaluations are essential for model engineering and its final deployment.
- For classification tasks, there are two types of evaluation metrics:
 - Discrete: use a threshold to binarize the predicted scores to 0/1. Compute the concordance between predicted labels and observed labels.
 - Continuous: keep predicted scores as continuous values. Compute the prediction consistence with observed labels over multiple thresholds.

See Also- Day 2 from the workshop: [Machine Learning with Python](#)


Discrete prediction evaluations

- Definition of TP, FP, TN, FN, f1 score.
- Confusion matrix is useful for visualizations, especially for multi-class predictions.



Continuous prediction evaluations

- Area under the Receiver operating curve (AUROC)
- Compares sensitivity (recall) to false positive rate (1-specificity) at various thresholds.
 - auROC = 1 (Perfect classifier)
 - auROC = 0.5 (Random classifier)



Model Save & Load

- Now that we have a preference over multiple models, it's essential to save a trained model and reload it for inference.

Keras

```
# saving a keras model is easy  
model_tf.save("trained_tf_model.h5")
```

```
# we can also load both the model and its weights from file  
model_tf_loaded = tf.keras.models.load_model("trained_tf_model.h5")  
model_tf_loaded.summary()
```

PyTorch

```
# by default, pytorch-lightning trainer will automatically save model checkpoints  
# see `lightning_logs/version_0/checkpoints`  
# we can also manually save pytorch model:  
trainer.save_checkpoint("trained_torch_model.ckpt")
```

```
# load from pytorch-lightning checkpoint  
model_torch_loaded = ModelTorch.load_from_checkpoint("trained_torch_model.ckpt")  
model_torch_loaded
```


Practice 3: Model Evaluations

*Run Section 3-5. Build a CNN Model using PyTorch/Lightning,
Model Inference and Evaluation and Model Saving and Loading*



- Understand PyTorch CNN model.
- Improve Model performance as evaluated by the metrics we covered.
- Compare different model's performance using different evaluation metrics, model architectures, and optimization configurations.

Summary & What's Next

- Understand how CNNs model genomic sequences
 - Build CNNs with Tensorflow and PyTorch
 - Model evaluations and tuning
-
- Tomorrow:
 - Framework-agnostic CNN models
 - Automatically find accurate CNNs in under 5min

Q&A

Write your questions in this [Google Doc](#).